

```
In [53]: import pandas as pd
import numpy as np
import imblearn
from imblearn.under_sampling import NearMiss
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV, train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.feature_selection import SequentialFeatureSelector
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.svm import SVC
from xgboost import XGBClassifier
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
import matplotlib.pyplot as plt
```

```
In [5]: #balancing the dataset here:
df = pd.read_csv('diabetes.csv')
undersample = NearMiss(version=1)
X = df.loc[:, df.columns != 'Diabetes_binary']
y = df.loc[:, df.columns == 'Diabetes_binary']
X, y = undersample.fit_resample(X, y)

#splitting the balanced dataset into train and testing samples
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state=42)
scaler = StandardScaler()
scaler.fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)

#putting the balanced datasets into individual dataframes for the train and test sets
df_undersampled_train = pd.DataFrame(X_train_scaled, columns = X.columns)
df_undersampled_train['Diabetes_binary'] = y_train
df_undersampled_train.head()

df_undersampled_test = pd.DataFrame(X_test_scaled, columns = X.columns)
df_undersampled_test['Diabetes_binary'] = y_test
df_undersampled_test.head()
```

```
Out[5]:
```

	HighBP	HighChol	CholCheck	BMI	Smoker	Stroke	HeartDiseaseorAttack	PhysA
0	-1.212894	0.876922	0.074482	-1.061978	1.158253	-0.225623	-0.384172	0.5
1	-1.212894	-1.140353	0.074482	0.377975	1.158253	-0.225623	-0.384172	0.5
2	0.824475	0.876922	0.074482	1.017954	-0.863369	-0.225623	-0.384172	0.5
3	-1.212894	0.876922	0.074482	0.377975	-0.863369	-0.225623	-0.384172	0.5
4	0.824475	-1.140353	0.074482	2.777896	-0.863369	-0.225623	-0.384172	0.5

5 rows × 22 columns

Features Choosing From LassoCV Selection Technique

```
In [51]: #Creating the dataframe that will contain the features that I selected through
X_selected_train = df_undersampled_train.loc[:, ['HighBP', 'BMI', 'Smoker', 'HeartDisease', 'GenHlth', 'MentHlth', 'PhysHlth']]

X_selected_test = df_undersampled_test.loc[:, ['HighBP', 'BMI', 'Smoker', 'HeartDisease', 'GenHlth', 'MentHlth', 'PhysHlth']]
```

Feature Selection From SFS Method

```
In [37]: sfs_selected_train = df_undersampled_train.iloc[:, [3, 5, 6, 7, 9, 13, 14, 15],]
sfs_selected_test = df_undersampled_test.iloc[:, [3, 5, 6, 7, 9, 13, 14, 15],]
```

XGBoost Model Creation and Parameter Tuning

```
In [38]: xgb_gs = GridSearchCV(estimator = XGBClassifier(),
                               param_grid={'objective': ['binary:logistic', 'binary:logit'],
                                             'learning_rate': [0.3, 0.6, 0.9],
                                             'max_depth': [6, 12, 18],
                                             'min_child_weight': [1, 5, 10],
                                             'subsamples': [0.5, 0.75, 1],
                                             'colsample_bytree': [0.5, 0.7, 1],
                                             'seed': [42]},
                               cv = 5,
                               verbose = 3,
                               n_jobs = -1)
xgb_gs.fit(X_selected_train, y_train.values.ravel())

print(xgb_gs.best_params_)
print(xgb_gs.best_score_)
```

Fitting 5 folds for each of 729 candidates, totalling 3645 fits
 {'colsample_bytree': 0.5, 'learning_rate': 0.3, 'max_depth': 6, 'min_child_weight': 1, 'objective': 'binary:logistic', 'seed': 42, 'subsamples': 0.5}
 0.8569840047721933

C:\Users\Felipe\anaconda3\lib\site-packages\xgboost\core.py:160: UserWarning:
 [16:02:00] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0750514818a16474a-1\xgboost\xgboost-ci-windows\src\learner.cc:742:
 Parameters: { "subsamples" } are not used.

warnings.warn(msg, UserWarning)

XGBoost Model Performance

Important to note what is a XGBoost model

- * It is an implementation of gradient boosting decision trees
- * Creates a series of models and combines them to make an overall model that is more accurate

Why to use XGBoost

- * It is a robust machine learning model
- * Easy to use
- * Fast on large datasets
- * There isn't much of a need for optimization but still encouraged
- * Offers built in regularization
 - * Meaning that it deals with non-zero entries and maintains its computational complexity like stochastic gradient descent
- * Easy to scale up on multicore machines or clusters
- * Uses disk based data structures instead of in-memory ones during computation

```
In [54]: xgb = XGBClassifier(colsample_bytree = 0.5, learning_rate = 0.3, max_depth=6,
                             objective='binary:hinge', seed=1, subsamples=0.5)
xgb.fit(X_selected_train, y_train.values.ravel())

xgb_train_pred = xgb.predict(X_selected_train)
xgb_train_score = accuracy_score(y_train.values.ravel(), xgb_train_pred)

xgb_test_pred = xgb.predict(X_selected_test)
xgb_test_score = accuracy_score(y_test.values.ravel(), xgb_test_pred)

print(f'Train Accuracy: {xgb_train_score}')
print(f'Classification Report for Train: \n{classification_report(y_train.values.ravel(), xgb_train_pred)}')
print(f'Confusion Matrix:\n{confusion_matrix(y_train.values.ravel(), xgb_train_pred)}')

print(f'Test Accuracy: {xgb_test_score}')
print(f'Classification Report for Test: \n{classification_report(y_test.values.ravel(), xgb_test_pred)}')
print(f'Confusion Matrix:\n{confusion_matrix(y_test.values.ravel(), xgb_test_pred)}')

disp = ConfusionMatrixDisplay(confusion_matrix(y_test.values.ravel(), xgb_test_pred))
disp.plot()
plt.show()
```

C:\Users\Felipe\anaconda3\lib\site-packages\xgboost\core.py:160: UserWarning:
[02:02:34] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscaling-group-i-0750514818a16474a-1\xgboost\xgboost-ci-windows\src\learner.cc:742:
Parameters: { "subsamples" } are not used.

```
warnings.warn(smsg, UserWarning)
```

Train Accuracy: 0.8570447013175976

Classification Report for Train:

	precision	recall	f1-score	support
0.0	0.81	0.94	0.87	24745
1.0	0.93	0.77	0.84	24739
accuracy			0.86	49484
macro avg	0.87	0.86	0.86	49484
weighted avg	0.87	0.86	0.86	49484

Confusion Matrix:

```
[[23271 1474]
 [ 5600 19139]]
```


Test Accuracy: 0.8518955111278763

Classification Report for Test:

	precision	recall	f1-score	support
0.0	0.80	0.94	0.86	10601
1.0	0.92	0.77	0.84	10607
accuracy			0.85	21208
macro avg	0.86	0.85	0.85	21208
weighted avg	0.86	0.85	0.85	21208

Confusion Matrix:

```
[[9926 675]
 [2466 8141]]
```



```
In [55]: xgb = XGBClassifier(colsample_bytree = 0.5, learning_rate = 0.3, max_depth=6, r
          objective='binary:hinge', seed=1, subsamples=0.5)
xgb.fit(sfs_selected_train, y_train.values.ravel())

xgb_train_pred = xgb.predict(sfs_selected_train)
xgb_train_score = accuracy_score(y_train.values.ravel(), xgb_train_pred)

xgb_test_pred = xgb.predict(sfs_selected_test)
xgb_test_score = accuracy_score(y_test.values.ravel(), xgb_test_pred)

print(f'Train Accuracy: {xgb_train_score}')
print(f'Classification Report for Train: \n{classification_report(y_train.values
print(f'Confusion Matrix:\n{confusion_matrix(y_train.values.ravel(), xgb_train

print(f'Test Accuracy: {xgb_test_score}')
print(f'Classification Report for Test: \n{classification_report(y_test.values
print(f'Confusion Matrix:\n{confusion_matrix(y_test.values.ravel(), xgb_test_p

disp = ConfusionMatrixDisplay(confusion_matrix(y_test.values.ravel(), xgb_test
disp.plot()
plt.show()
```

C:\Users\Felipe\anaconda3\lib\site-packages\xgboost\core.py:160: UserWarning:
 [02:02:55] WARNING: C:\buildkite-agent\builds\buildkite-windows-cpu-autoscali
 ng-group-i-0750514818a16474a-1\xgboost\xgboost-ci-windows\src\learner.cc:742:
 Parameters: { "subsamples" } are not used.

```
warnings.warn(smsg, UserWarning)
```

Train Accuracy: 0.8809514186403686

Classification Report for Train:

	precision	recall	f1-score	support
0.0	0.84	0.95	0.89	24745
1.0	0.94	0.82	0.87	24739
accuracy			0.88	49484
macro avg	0.89	0.88	0.88	49484
weighted avg	0.89	0.88	0.88	49484

Confusion Matrix:

```
[[23412 1333]
 [ 4558 20181]]
```


Test Accuracy: 0.8741512636740852

Classification Report for Test:

	precision	recall	f1-score	support
0.0	0.83	0.94	0.88	10601
1.0	0.93	0.81	0.87	10607
accuracy			0.87	21208
macro avg	0.88	0.87	0.87	21208
weighted avg	0.88	0.87	0.87	21208

Confusion Matrix:

```
[[9953 648]
 [2021 8586]]
```

Extreme Machine Learning Model

- While it is a feed forward neural network it has a few differences
 - First, is that it does not use back propogation
 - Second, the weights for the hidden layers are randomly generated
 - Lastly, the output is linear rather than sigmoidal

```
In [41]: input_length = X_selected_train.shape[1]
hidden_units = 1000

win = np.random.normal(size = [input_length, hidden_units])

def input_to_hidden(x):
    a = np.dot(x, win)
    a = np.maximum(a, 0, a)
    return a

x_h_v = input_to_hidden(X_selected_train)
x_h_t = np.transpose(x_h_v)
w_out = np.dot(np.linalg.inv(np.dot(x_h_t, x_h_v)), np.dot(x_h_t, y_train))

def predict(x):
    x = input_to_hidden(x)
    y = np.dot(x, w_out)
    return y

extreme_pred = predict(X_selected_test)
num_correct = 0
total = extreme_pred.shape[0]
for i in range(total):
    predicted = np.argmax(extreme_pred[i])
    test = np.argmax(y_test.values.ravel()[i])
    num_correct = num_correct + (1 if predicted == test else 0)

print('Accuracy of test set: {:.f}'.format(num_correct/total))

extreme_pred = predict(X_selected_train)
num_correct = 0
total = extreme_pred.shape[0]
for i in range(total):
    predicted = np.argmax(extreme_pred[i])
    train = np.argmax(y_train.values.ravel()[i])
    num_correct = num_correct + (1 if predicted == train else 0)

print('Accuracy of train set: {:.f}'.format(num_correct/total))
```

Accuracy of test set: 1.000000

Accuracy of train set: 1.000000

```

In [42]: input_length = sfs_selected_train.shape[1]
hidden_units = 1000

win = np.random.normal(size = [input_length, hidden_units])

def input_to_hidden(x):
    a = np.dot(x, win)
    a = np.maximum(a, 0, a)
    return a

x_h_v = input_to_hidden(sfs_selected_train)
x_h_t = np.transpose(x_h_v)
w_out = np.dot(np.linalg.inv(np.dot(x_h_t, x_h_v)), np.dot(x_h_t, y_train))

def predict(x):
    x = input_to_hidden(x)
    y = np.dot(x, w_out)
    return y

extreme_pred = predict(sfs_selected_test)
num_correct = 0
total = extreme_pred.shape[0]
for i in range(total):
    predicted = np.argmax(extreme_pred[i])
    test = np.argmax(y_test.values.ravel()[i])
    num_correct = num_correct + (1 if predicted == test else 0)

print('Accuracy of test set: {:.f}'.format(num_correct/total))

extreme_pred = predict(sfs_selected_train)
num_correct = 0
total = extreme_pred.shape[0]
for i in range(total):
    predicted = np.argmax(extreme_pred[i])
    train = np.argmax(y_train.values.ravel()[i])
    num_correct = num_correct + (1 if predicted == train else 0)

print('Accuracy of train set: {:.f}'.format(num_correct/total))

```

Accuracy of test set: 1.000000

Accuracy of train set: 1.000000

Basic Deep Learning Model

- Basic feed forward neural network
 - This neural network has 2 hidden layers, 1 input and 1 output layer

```
In [72]: nn = Sequential()
nn.add(Dense(20, input_shape=(9,), activation='relu'))
nn.add(Dense(40, activation='relu'))
nn.add(Dense(10, activation = 'relu'))
nn.add(Dense(1, activation='sigmoid'))

nn.compile(loss = 'binary_crossentropy', optimizer='adam', metrics='accuracy')
nn.fit(X_selected_train, y_train.values.ravel(), epochs=50, batch_size=10)

_, train_accuracy = nn.evaluate(X_selected_train, y_train)
_, test_accuracy = nn.evaluate(X_selected_test, y_test)

print(f'Train Accuracy: {train_accuracy}')
print(f'Test Accuracy: {test_accuracy}')
```

```
Epoch 1/50
4949/4949 [=====] - 4s 722us/step - loss: 0.3511 - a
ccuracy: 0.8416
Epoch 2/50
4949/4949 [=====] - 4s 734us/step - loss: 0.3299 - a
ccuracy: 0.8520
Epoch 3/50
4949/4949 [=====] - 4s 763us/step - loss: 0.3263 - a
ccuracy: 0.8543
Epoch 4/50
4949/4949 [=====] - 4s 757us/step - loss: 0.3235 - a
ccuracy: 0.8544
Epoch 5/50
4949/4949 [=====] - 4s 769us/step - loss: 0.3231 - a
ccuracy: 0.8538
Epoch 6/50
4949/4949 [=====] - 4s 734us/step - loss: 0.3225 - a
ccuracy: 0.8542
Epoch 7/50
4949/4949 [=====] - 4s 734us/step - loss: 0.3216 - a
ccuracy: 0.8554
Epoch 8/50
4949/4949 [=====] - 4s 736us/step - loss: 0.3209 - a
ccuracy: 0.8557
Epoch 9/50
4949/4949 [=====] - 4s 733us/step - loss: 0.3204 - a
ccuracy: 0.8564
Epoch 10/50
4949/4949 [=====] - 4s 726us/step - loss: 0.3199 - a
ccuracy: 0.8554
Epoch 11/50
4949/4949 [=====] - 4s 723us/step - loss: 0.3199 - a
ccuracy: 0.8562
Epoch 12/50
4949/4949 [=====] - 4s 730us/step - loss: 0.3199 - a
ccuracy: 0.8562
Epoch 13/50
4949/4949 [=====] - 4s 729us/step - loss: 0.3193 - a
ccuracy: 0.8559
Epoch 14/50
4949/4949 [=====] - 4s 733us/step - loss: 0.3196 - a
ccuracy: 0.8556
Epoch 15/50
4949/4949 [=====] - 4s 740us/step - loss: 0.3192 - a
ccuracy: 0.8567
Epoch 16/50
4949/4949 [=====] - 4s 736us/step - loss: 0.3191 - a
ccuracy: 0.8561
Epoch 17/50
4949/4949 [=====] - 4s 732us/step - loss: 0.3184 - a
ccuracy: 0.8568
Epoch 18/50
4949/4949 [=====] - 4s 728us/step - loss: 0.3184 - a
ccuracy: 0.8566
Epoch 19/50
4949/4949 [=====] - 4s 728us/step - loss: 0.3186 - a
ccuracy: 0.8571
```

```
Epoch 20/50
4949/4949 [=====] - 4s 733us/step - loss: 0.3183 - a
ccuracy: 0.8570
Epoch 21/50
4949/4949 [=====] - 4s 735us/step - loss: 0.3186 - a
ccuracy: 0.8570
Epoch 22/50
4949/4949 [=====] - 4s 731us/step - loss: 0.3182 - a
ccuracy: 0.8572
Epoch 23/50
4949/4949 [=====] - 4s 738us/step - loss: 0.3178 - a
ccuracy: 0.8566
Epoch 24/50
4949/4949 [=====] - 4s 745us/step - loss: 0.3176 - a
ccuracy: 0.8567
Epoch 25/50
4949/4949 [=====] - 4s 729us/step - loss: 0.3178 - a
ccuracy: 0.8575
Epoch 26/50
4949/4949 [=====] - 4s 740us/step - loss: 0.3176 - a
ccuracy: 0.8573
Epoch 27/50
4949/4949 [=====] - 4s 731us/step - loss: 0.3175 - a
ccuracy: 0.8567
Epoch 28/50
4949/4949 [=====] - 4s 725us/step - loss: 0.3176 - a
ccuracy: 0.8577
Epoch 29/50
4949/4949 [=====] - 4s 747us/step - loss: 0.3175 - a
ccuracy: 0.8576
Epoch 30/50
4949/4949 [=====] - 4s 738us/step - loss: 0.3174 - a
ccuracy: 0.8567
Epoch 31/50
4949/4949 [=====] - 4s 733us/step - loss: 0.3173 - a
ccuracy: 0.8572
Epoch 32/50
4949/4949 [=====] - 4s 744us/step - loss: 0.3173 - a
ccuracy: 0.8569
Epoch 33/50
4949/4949 [=====] - 4s 743us/step - loss: 0.3173 - a
ccuracy: 0.8568
Epoch 34/50
4949/4949 [=====] - 4s 722us/step - loss: 0.3175 - a
ccuracy: 0.8572
Epoch 35/50
4949/4949 [=====] - 4s 729us/step - loss: 0.3169 - a
ccuracy: 0.8571
Epoch 36/50
4949/4949 [=====] - 4s 738us/step - loss: 0.3172 - a
ccuracy: 0.8579
Epoch 37/50
4949/4949 [=====] - 4s 720us/step - loss: 0.3170 - a
ccuracy: 0.8574
Epoch 38/50
4949/4949 [=====] - 4s 736us/step - loss: 0.3169 - a
ccuracy: 0.8573
```


```
Epoch 39/50
4949/4949 [=====] - 4s 732us/step - loss: 0.3167 - a
ccuracy: 0.8579
Epoch 40/50
4949/4949 [=====] - 4s 764us/step - loss: 0.3168 - a
ccuracy: 0.8573
Epoch 41/50
4949/4949 [=====] - 4s 738us/step - loss: 0.3165 - a
ccuracy: 0.8574
Epoch 42/50
4949/4949 [=====] - 4s 721us/step - loss: 0.3169 - a
ccuracy: 0.8573
Epoch 43/50
4949/4949 [=====] - 4s 736us/step - loss: 0.3165 - a
ccuracy: 0.8574
Epoch 44/50
4949/4949 [=====] - 4s 730us/step - loss: 0.3163 - a
ccuracy: 0.8571
Epoch 45/50
4949/4949 [=====] - 4s 759us/step - loss: 0.3166 - a
ccuracy: 0.8575
Epoch 46/50
4949/4949 [=====] - 4s 761us/step - loss: 0.3164 - a
ccuracy: 0.8582
Epoch 47/50
4949/4949 [=====] - 4s 757us/step - loss: 0.3163 - a
ccuracy: 0.8576
Epoch 48/50
4949/4949 [=====] - 4s 735us/step - loss: 0.3163 - a
ccuracy: 0.8573
Epoch 49/50
4949/4949 [=====] - 4s 754us/step - loss: 0.3162 - a
ccuracy: 0.8572
Epoch 50/50
4949/4949 [=====] - 4s 765us/step - loss: 0.3162 - a
ccuracy: 0.8572
1547/1547 [=====] - 1s 632us/step - loss: 0.3139 - a
ccuracy: 0.8587
663/663 [=====] - 0s 613us/step - loss: 0.3279 - acc
uracy: 0.8535
Train Accuracy: 0.8587422370910645
Test Accuracy: 0.8534515500068665
```

```
In [73]: nn_test_pred = nn.predict(X_selected_test)
nn_test_pred = np.round(nn_test_pred)

cm = confusion_matrix(y_test, nn_test_pred)

disp = ConfusionMatrixDisplay(confusion_matrix(y_test.values.ravel(), nn_test_
disp.plot()
plt.show()

print(f'Test Accuracy: {test_accuracy}')
print(f'Classification Report for Test: \n{classification_report(y_test.values
663/663 [=====] - 0s 621us/step
```



Test Accuracy: 0.8534515500068665

Classification Report for Test:

	precision	recall	f1-score	support
0.0	0.80	0.93	0.86	10601
1.0	0.92	0.77	0.84	10607
accuracy			0.85	21208
macro avg	0.86	0.85	0.85	21208
weighted avg	0.86	0.85	0.85	21208


```
In [68]: nn = Sequential()
nn.add(Dense(20, input_shape=(9,), activation='relu'))
nn.add(Dense(40, activation='relu'))
nn.add(Dense(10, activation = 'relu'))
nn.add(Dense(1, activation='sigmoid'))

nn.compile(loss = 'binary_crossentropy', optimizer='adam', metrics='accuracy')
nn.fit(sfs_selected_train, y_train.values.ravel(), epochs=50, batch_size=10)

_, train_accuracy = nn.evaluate(sfs_selected_train, y_train)
_, test_accuracy = nn.evaluate(sfs_selected_test, y_test)

print(f'Train Accuracy: {train_accuracy}')
print(f'Test Accuracy: {test_accuracy}')
```

```
Epoch 1/50
4949/4949 [=====] - 4s 729us/step - loss: 0.3124 - a
ccuracy: 0.8634
Epoch 2/50
4949/4949 [=====] - 4s 728us/step - loss: 0.2928 - a
ccuracy: 0.8752
Epoch 3/50
4949/4949 [=====] - 4s 730us/step - loss: 0.2896 - a
ccuracy: 0.8757
Epoch 4/50
4949/4949 [=====] - 4s 743us/step - loss: 0.2883 - a
ccuracy: 0.8760
Epoch 5/50
4949/4949 [=====] - 4s 735us/step - loss: 0.2871 - a
ccuracy: 0.8759
Epoch 6/50
4949/4949 [=====] - 4s 745us/step - loss: 0.2861 - a
ccuracy: 0.8761
Epoch 7/50
4949/4949 [=====] - 4s 748us/step - loss: 0.2857 - a
ccuracy: 0.8776
Epoch 8/50
4949/4949 [=====] - 4s 738us/step - loss: 0.2850 - a
ccuracy: 0.8778
Epoch 9/50
4949/4949 [=====] - 4s 724us/step - loss: 0.2850 - a
ccuracy: 0.8775
Epoch 10/50
4949/4949 [=====] - 4s 725us/step - loss: 0.2839 - a
ccuracy: 0.8779
Epoch 11/50
4949/4949 [=====] - 4s 724us/step - loss: 0.2843 - a
ccuracy: 0.8783
Epoch 12/50
4949/4949 [=====] - 4s 728us/step - loss: 0.2839 - a
ccuracy: 0.8775
Epoch 13/50
4949/4949 [=====] - 4s 730us/step - loss: 0.2834 - a
ccuracy: 0.8772
Epoch 14/50
4949/4949 [=====] - 4s 728us/step - loss: 0.2836 - a
ccuracy: 0.8780
Epoch 15/50
4949/4949 [=====] - 4s 730us/step - loss: 0.2832 - a
ccuracy: 0.8775
Epoch 16/50
4949/4949 [=====] - 4s 727us/step - loss: 0.2829 - a
ccuracy: 0.8786
Epoch 17/50
4949/4949 [=====] - 4s 725us/step - loss: 0.2828 - a
ccuracy: 0.8779
Epoch 18/50
4949/4949 [=====] - 4s 723us/step - loss: 0.2822 - a
ccuracy: 0.8785
Epoch 19/50
4949/4949 [=====] - 4s 725us/step - loss: 0.2828 - a
ccuracy: 0.8782
```

```
Epoch 20/50
4949/4949 [=====] - 4s 724us/step - loss: 0.2821 - a
ccuracy: 0.8786
Epoch 21/50
4949/4949 [=====] - 4s 751us/step - loss: 0.2824 - a
ccuracy: 0.8783
Epoch 22/50
4949/4949 [=====] - 4s 730us/step - loss: 0.2817 - a
ccuracy: 0.8788
Epoch 23/50
4949/4949 [=====] - 4s 733us/step - loss: 0.2824 - a
ccuracy: 0.8785
Epoch 24/50
4949/4949 [=====] - 4s 738us/step - loss: 0.2818 - a
ccuracy: 0.8777
Epoch 25/50
4949/4949 [=====] - 4s 737us/step - loss: 0.2818 - a
ccuracy: 0.8779
Epoch 26/50
4949/4949 [=====] - 4s 752us/step - loss: 0.2810 - a
ccuracy: 0.8796
Epoch 27/50
4949/4949 [=====] - 4s 733us/step - loss: 0.2814 - a
ccuracy: 0.8787
Epoch 28/50
4949/4949 [=====] - 4s 722us/step - loss: 0.2815 - a
ccuracy: 0.8785
Epoch 29/50
4949/4949 [=====] - 4s 721us/step - loss: 0.2810 - a
ccuracy: 0.8779
Epoch 30/50
4949/4949 [=====] - 4s 729us/step - loss: 0.2810 - a
ccuracy: 0.8785
Epoch 31/50
4949/4949 [=====] - 4s 727us/step - loss: 0.2810 - a
ccuracy: 0.8783
Epoch 32/50
4949/4949 [=====] - 4s 727us/step - loss: 0.2810 - a
ccuracy: 0.8788
Epoch 33/50
4949/4949 [=====] - 4s 726us/step - loss: 0.2805 - a
ccuracy: 0.8787
Epoch 34/50
4949/4949 [=====] - 4s 729us/step - loss: 0.2805 - a
ccuracy: 0.8796
Epoch 35/50
4949/4949 [=====] - 4s 736us/step - loss: 0.2806 - a
ccuracy: 0.8789
Epoch 36/50
4949/4949 [=====] - 4s 735us/step - loss: 0.2804 - a
ccuracy: 0.8785
Epoch 37/50
4949/4949 [=====] - 4s 737us/step - loss: 0.2803 - a
ccuracy: 0.8784
Epoch 38/50
4949/4949 [=====] - 4s 732us/step - loss: 0.2809 - a
ccuracy: 0.8790
```


```
Epoch 39/50
4949/4949 [=====] - 4s 735us/step - loss: 0.2805 - a
ccuracy: 0.8793
Epoch 40/50
4949/4949 [=====] - 4s 735us/step - loss: 0.2800 - a
ccuracy: 0.8794
Epoch 41/50
4949/4949 [=====] - 4s 732us/step - loss: 0.2804 - a
ccuracy: 0.8797
Epoch 42/50
4949/4949 [=====] - 4s 730us/step - loss: 0.2801 - a
ccuracy: 0.8788
Epoch 43/50
4949/4949 [=====] - 4s 741us/step - loss: 0.2800 - a
ccuracy: 0.8797
Epoch 44/50
4949/4949 [=====] - 4s 737us/step - loss: 0.2801 - a
ccuracy: 0.8797
Epoch 45/50
4949/4949 [=====] - 4s 724us/step - loss: 0.2797 - a
ccuracy: 0.8798
Epoch 46/50
4949/4949 [=====] - 4s 734us/step - loss: 0.2797 - a
ccuracy: 0.8787
Epoch 47/50
4949/4949 [=====] - 4s 730us/step - loss: 0.2799 - a
ccuracy: 0.8800
Epoch 48/50
4949/4949 [=====] - 4s 730us/step - loss: 0.2798 - a
ccuracy: 0.8797
Epoch 49/50
4949/4949 [=====] - 4s 748us/step - loss: 0.2795 - a
ccuracy: 0.8797
Epoch 50/50
4949/4949 [=====] - 4s 728us/step - loss: 0.2796 - a
ccuracy: 0.8795
1547/1547 [=====] - 1s 592us/step - loss: 0.2776 - a
ccuracy: 0.8804
663/663 [=====] - 0s 589us/step - loss: 0.3010 - acc
uracy: 0.8747
Train Accuracy: 0.880446195602417
Test Accuracy: 0.8746699094772339
```

```
In [70]: nn_test_pred = nn.predict(sfs_selected_test)
nn_test_pred = np.round(nn_test_pred)

cm = confusion_matrix(y_test, nn_test_pred)

disp = ConfusionMatrixDisplay(confusion_matrix(y_test.values.ravel(), nn_test_
disp.plot()
plt.show()

print(f'Test Accuracy: {test_accuracy}')
print(f'Classification Report for Test: \n{classification_report(y_test.values
663/663 [=====] - 0s 544us/step
```



Test Accuracy: 0.8746699094772339

Classification Report for Test:

	precision	recall	f1-score	support
0.0	0.84	0.93	0.88	10601
1.0	0.92	0.82	0.87	10607
accuracy			0.87	21208
macro avg	0.88	0.87	0.87	21208
weighted avg	0.88	0.87	0.87	21208

Ensemble Model

- * This model combines a n number of models into one and decides final prediction through various methods
- * This model uses the voting method
 - * With hard voting so the predictions are based on majority vote

```
In [66]: model_1 = RandomForestClassifier(criterion = 'entropy', max_features = 'sqrt',
model_2 = SVC(C = 5, degree = 3, gamma = 'auto', kernel = 'rbf')
model_3 = LogisticRegression(C = 8, l1_ratio = 0, penalty = 'l1', random_state=

ensemble_model = VotingClassifier(estimators = [('rf', model_1), ('svm', model_
                                voting = 'hard',
                                n_jobs = -1,
                                verbose = True)
ensemble_model.fit(X_selected_train, y_train.values.ravel())
ensemble_pred_train = ensemble_model.predict(X_selected_train)
ensemble_pred_test = ensemble_model.predict(X_selected_test)

ensemble_train_score = accuracy_score(y_train.values.ravel(), ensemble_pred_train)
ensemble_test_score = accuracy_score(y_test.values.ravel(), ensemble_pred_test)

print(f'Train Accuracy: {ensemble_train_score}')
print(f'Classification Report for Train: \n{classification_report(y_train.values
print(f'Test Accuracy: {ensemble_test_score}')
print(f'Classification Report for Test: \n{classification_report(y_test.values
disp = ConfusionMatrixDisplay(confusion_matrix(y_test.values.ravel(), ensemble
disp.plot()
plt.show()
```

Train Accuracy: 0.8555290598981489


Classification Report for Train:

	precision	recall	f1-score	support
0.0	0.81	0.93	0.87	24745
1.0	0.92	0.78	0.84	24739
accuracy			0.86	49484
macro avg	0.86	0.86	0.85	49484
weighted avg	0.86	0.86	0.85	49484

Test Accuracy: 0.8515654470011317

Classification Report for Test:

	precision	recall	f1-score	support
0.0	0.81	0.93	0.86	10601
1.0	0.91	0.78	0.84	10607
accuracy			0.85	21208
macro avg	0.86	0.85	0.85	21208
weighted avg	0.86	0.85	0.85	21208




```
In [67]: model_1 = RandomForestClassifier(criterion = 'entropy', max_features = 'sqrt',
model_2 = SVC(C = 5, degree = 3, gamma = 'auto', kernel = 'rbf')
model_3 = LogisticRegression(C = 8, l1_ratio = 0, penalty = 'l1', random_state=

ensemble_model = VotingClassifier(estimators = [('rf', model_1), ('svm', model_
                                voting = 'hard',
                                n_jobs = -1,
                                verbose = True)
ensemble_model.fit(sfs_selected_train, y_train.values.ravel())
ensemble_pred_train = ensemble_model.predict(sfs_selected_train)
ensemble_pred_test = ensemble_model.predict(sfs_selected_test)

ensemble_train_score = accuracy_score(y_train.values.ravel(), ensemble_pred_train)
ensemble_test_score = accuracy_score(y_test.values.ravel(), ensemble_pred_test)

print(f'Train Accuracy: {ensemble_train_score}')
print(f'Classification Report for Train: \n{classification_report(y_train.values
print(f'Test Accuracy: {ensemble_test_score}')
print(f'Classification Report for Test: \n{classification_report(y_test.values

disp = ConfusionMatrixDisplay(confusion_matrix(y_test.values.ravel(), ensemble
disp.plot()
plt.show()
```

Train Accuracy: 0.8776574246221001


Classification Report for Train:

	precision	recall	f1-score	support
0.0	0.84	0.94	0.88	24745
1.0	0.93	0.82	0.87	24739
accuracy			0.88	49484
macro avg	0.88	0.88	0.88	49484
weighted avg	0.88	0.88	0.88	49484

Test Accuracy: 0.872736703130894

Classification Report for Test:

	precision	recall	f1-score	support
0.0	0.83	0.93	0.88	10601
1.0	0.92	0.82	0.87	10607
accuracy			0.87	21208
macro avg	0.88	0.87	0.87	21208
weighted avg	0.88	0.87	0.87	21208



Once Again The Wrapper Method Seems To Offer The Best Results